

PROFESSIONAL SUMMARY

- Database engineer with 5+ years designing, developing, and maintaining large-scale relational database systems — with a track record of high-availability SQL server operations, stored procedure development, schema optimisation, and measurable performance gains.
- Proven ability to debug, optimise, and support production database environments: delivered a 97% reduction in database refresh time and maintained MySQL servers at 99% availability under high-throughput production workloads.
- Hands-on Linux server experience — managing and monitoring database services, writing Bash and Python automation scripts, analysing logs, and resolving service issues in production environments.
- Azure cloud experience: designed and optimised Azure VM-hosted database infrastructure, reducing cloud costs 30% through database server right-sizing and resource utilisation improvements.

TECHNICAL SKILLS

Database Technologies: SQL Server (T-SQL), MySQL, PostgreSQL, MongoDB — scripting, indices, views, stored procedures, bulk insert operations, materialized views

Database Design: Logical and physical data modeling, schema design, query optimisation, data integrity, ORM patterns, ETL pipeline design

Azure: Azure VM (database server hosting, scaling, right-sizing), Azure Service Bus (data ingestion pipelines), Azure Blob Storage

Linux / Server: Database service management, process monitoring, log analysis (EFK Stack), Bash and Python scripting for automation, Docker, Kubernetes

DevOps: GitHub Actions, CI/CD Pipelines, AWS (EC2), Git, Agile / Scrum

Languages: Python (asyncio), SQL, Bash, TypeScript, Java

EXPERIENCE

Lead Developer | 08/2025 – 12/2025 | ABC Academy of Music | MS Industry Capstone | Phoenix, AZ

Led database architecture and backend design for a production platform, managing a team of 5 in an Agile environment.

- Designed the PostgreSQL relational database schema from scratch — translating business requirements into logical and physical data models, defining table structures, relationships, and indices to support a Java Spring Boot API.
- Deployed and maintained the database-backed application on AWS EC2 Linux servers, configuring CI/CD pipelines for automated deployments and achieving 99% platform availability.
- Documented database design decisions and schema changes as part of rigorous Agile delivery standards, maintaining clear records for team handoff and post-implementation support.

Lead Developer | 04/2023 – 12/2023 | Tecnotree | Bengaluru, KA

Awarded the "Spotlight Top Performer" cash award for leading high-quality, reliable feature delivery in a fast-paced Agile environment.

Senior Software Engineer | 07/2022 – 03/2023 | PangeaTech | Bengaluru, KA

Founding engineer on a production ML platform — owned database-backed pipeline architecture, Azure cloud infrastructure, and Linux server operations from inception through production.

- **Azure database infrastructure:** Designed and managed database services hosted on Azure VMs — optimising VM sizing, scaling Node.js database-serving processes across CPU cores using PM2 cluster mode, and reducing Azure cloud costs by 30% through better server resource utilisation.
- **Linux server operations:** Deployed and maintained all platform services on Linux servers — managing database processes, configuring services, monitoring system health, and resolving production incidents using shell scripting and system tools.
- **Real-time database monitoring:** Replaced manual server log inspection with a centralised EFK stack (Elasticsearch, Fluentd, Kibana) on Linux — reducing database pipeline error identification from 4 hours to real time, enabling faster post-implementation debugging and support.
- Scaled database-intensive Python data processing pipelines across Kubernetes nodes on Azure, reducing processing time by 80% (4 hours → 45 minutes) and relieving load on the production database layer.
- Authored 50+ code reviews and mentored junior engineers on SQL best practices, database design, and query performance standards.

Software Engineer | 11/2021 – 07/2022 | PangeaTech | Bengaluru, KA

Primary database engineer — sole owner of SQL server maintenance, stored procedure development, schema design, data ingestion pipelines, and 24/7 high-availability operations on Linux.

- **Maintained MySQL servers at 99% availability** on Linux — actively monitoring database service health, debugging and optimising SQL queries, maintaining stored procedures, and resolving production incidents to ensure continuous uptime.
- **97% database performance gain:** Engineered a 97% reduction in database refresh duration (8 hours → 15 minutes) by redesigning synchronous Python ETL scripts into fully asynchronous Python (asyncio), directly eliminating a critical production bottleneck.
- **High-throughput data ingestion:** Designed and executed bulk SQL insert pipelines — batching write requests and controlling throughput via API rate limiting to prevent out-of-memory exceptions during large-scale data collection jobs.

- Designed and maintained 50+ SQL materialized views powering real-time client dashboards — enforcing data integrity, index strategies, and schema consistency across a high-read production environment.
- **Linux server administration:** Managed database and application services on Linux servers — writing Bash and Python scripts to automate routine database maintenance, monitoring logs, managing processes, and troubleshooting service failures.
- Containerised database-backed application services with Docker and automated deployments via GitHub Actions, reducing deployment time by 85% (27 min → 4 min) and eliminating deployment-related database downtime.

Software Engineer | 07/2019 – 04/2021 | Tenable (a JPMorgan Chase company) | Pune, MH

Worked on a high-traffic platform serving millions of customers — keeping database security, data integrity, and long-term maintainability at the core of all engineering decisions.

ADDITIONAL EXPERIENCE

Music Recommendation & Classification | Personal Project

- Designed a relational + vector database system using PostgreSQL + pgvector — modelling schema for embedding-based semantic search, demonstrating initiative in applying emerging database technologies to solve real retrieval problems.

Engineering Tutor | ASU | 08/2024 – 12/2025

- Delivered 100+ tutoring sessions focused on Databases, Data Structures and Algorithms, and Operating Systems — reinforcing deep command of SQL, database internals, and system-level concepts.

EDUCATION

Master of Science in Software Engineering, GPA: 3.87 / 4.0 | Dec 2025 | Arizona State University | Mesa, AZ

Coursework: System Design, Data Structures and Algorithms, User Experience, Agile Software Development

Bachelor of Science in Computer Science | Aug 2019 | REVA University | Bengaluru, KA